

Table S3. Trained models at $g = 5$ km in the four Bauru cells (five draws)

A trained model’s valid-node MAE drifts with the split draw. Split draws 101–105 in the four Bauru cells, $g = 5$ km, $b = 2$ km, training seed 42, training budget of eight epochs. Valid-node mean absolute error of received power, in decibels, mean over draws (between-draw standard deviation) for the graph regressor (GNN) and the perceptron (MLP); bottom row: standard deviation of the per-cell means. Over the five draws, the between-draw standard deviation of the valid-node MAE is 5.2 to 21.6 times (GNN) and 4.7 to 32.0 times (MLP) the same-seed repeat difference (0.132 dB, the largest difference between two training repeats with the same seed and partition, measured at 5 km). In Bauru Q1 and Bauru Q3, where five training seeds were also run at split seed 42, it is 11.7 (GNN) and 38.2 (MLP) times the standard deviation across the five training seeds in Bauru Q1, and 2.8 (GNN) and 50.6 (MLP) times in Bauru Q3.

Cell	GNN valid MAE ↓ (dB): mean (SD over draws)	MLP valid MAE ↓ (dB): mean (SD over draws)
Bauru Q1	1.80 (1.25)	3.81 (1.71)
Bauru Q2	2.90 (2.85)	6.03 (4.22)
Bauru Q3	1.61 (0.98)	3.09 (2.93)
Bauru Q4	1.11 (0.69)	1.57 (0.62)
SD across the four cells of the per-cell means	0.76	1.85